

Measures of Distance in Data Mining

How similar are two data points?

Why Measure Distance?

Euclidean Distance

Manhattan Distance

Jaccard Index

Minkowski Distance

Cosine Similarity

Some More Distance Measures

Why Measure Distance?

Why Measure Distance? (Theory)

Distance measures quantify the similarity or dissimilarity between data points. They are foundational in:

- Clustering (e.g., K-Means, Hierarchical)
- Classification (e.g., K-Nearest Neighbors)
- Anomaly detection
- Information retrieval and recommender systems

Choosing the right distance measure affects model accuracy and interpretability.

Euclidean Distance

Euclidean Distance (Theory)

Euclidean distance is the most common measure of straight-line distance in continuous space.

- Sensitive to magnitude and scale
- Assumes continuous, real-valued attributes

Euclidean Distance (Math)

Given two points $A = (x_1, x_2, \dots, x_n)$ and $B = (y_1, y_2, \dots, y_n)$, the Euclidean distance is:

$$d(A, B) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Euclidean Distance (Example)

Given $A = (2, 4)$, $B = (5, 8)$:

$$d(A, B) = \sqrt{(2 - 5)^2 + (4 - 8)^2} = \sqrt{(-3)^2 + (-4)^2} = \sqrt{9 + 16} = \sqrt{25}$$

Euclidean Distance (Use Case)

Use Case: K-Means Clustering

- Used to assign points to the nearest centroid
- Ideal when clusters are spherical and scales are normalized

Manhattan Distance

Manhattan Distance (Theory)

Also called ****Taxicab**** or ****City Block Distance****:

- Measures distance by summing absolute differences
- Better for high-dimensional, sparse data

Manhattan Distance (Math)

$$d(A, B) = \sum_{i=1}^n |x_i - y_i|$$

Where $A = (x_1, \dots, x_n)$, $B = (y_1, \dots, y_n)$

Manhattan Distance (Example)

Given $A = (2, 4)$, $B = (5, 8)$:

$$d(A, B) = |2 - 5| + |4 - 8| = 3 + 4 = 7$$

Manhattan Distance (Use Case)

Use Case: L1-Regularized models (e.g., Lasso Regression)

- Promotes sparsity
- Preferred when feature differences are linear or additive

Jaccard Index

Jaccard Index (Theory)

Jaccard Similarity measures overlap between sets.

- Works with binary or categorical data
- Good for market basket analysis, document similarity

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

$$d(A, B) = 1 - J(A, B)$$

Jaccard Index (Example)

Let $A = \{1, 2, 3, 5\}$, $B = \{2, 3, 4, 6\}$

$$A \cap B = \{2, 3\}, \quad A \cup B = \{1, 2, 3, 4, 5, 6\}$$

$$J(A, B) = \frac{2}{6} = 0.33, \quad d(A, B) = 1 - 0.33 = 0.67$$

Jaccard Index (Use Case)

Use Case: Recommender Systems

- Used to compare user interests (e.g., liked items)
- Suitable for sparse, binary user-item matrices

Minkowski Distance

Minkowski Distance (Theory)

Generalized distance metric that includes:

- Euclidean distance when $p = 2$
- Manhattan distance when $p = 1$

Minkowski Distance (Math)

$$d(A, B) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

Where $p \in \mathbb{R}$, $p \geq 1$

Minkowski Distance (Example)

Given $A = (1, 2)$, $B = (4, 6)$, with $p = 3$:

$$d(A, B) = (|1 - 4|^3 + |2 - 6|^3)^{1/3} = (27 + 64)^{1/3} = (91)^{1/3} \approx 4.481$$

Minkowski Distance (Use Case)

Use Case: Customizable distance metric in KNN

- Choose p based on desired sensitivity
- Flexible for tuning similarity in various data distributions

Cosine Similarity

Cosine Similarity (Theory)

Cosine Similarity measures the cosine of the angle between two vectors.

- Focuses on orientation, not magnitude
- Ideal for high-dimensional, sparse data like text

Cosine Similarity (Math)

$$\text{Cosine}(A, B) = \frac{A \cdot B}{\|A\| \cdot \|B\|} = \frac{\sum_{i=1}^n x_i y_i}{\sqrt{\sum x_i^2} \cdot \sqrt{\sum y_i^2}}$$

Cosine Similarity (Example)

Example 1: Similar Vectors

Let $A_1 = (1, 2, 3)$, $B_1 = (2, 4, 6)$

$$\cos(\theta) = \frac{1 \cdot 2 + 2 \cdot 4 + 3 \cdot 6}{\sqrt{1^2 + 2^2 + 3^2} \cdot \sqrt{2^2 + 4^2 + 6^2}} = \frac{28}{\sqrt{14} \cdot \sqrt{56}} = 1$$

Perfect similarity — linearly dependent or colinear (same direction)

Example 2: Dissimilar Vectors

Let $A_2 = (1, 0)$, $B_2 = (0, 1)$

$$\cos(\theta) = \frac{1 \cdot 0 + 0 \cdot 1}{\sqrt{1^2 + 0^2} \cdot \sqrt{0^2 + 1^2}} = 0$$

No similarity — orthogonal vectors.

Cosine Similarity (Use Case)

Use Case: Document Similarity in NLP

- Used to compare TF-IDF or word embeddings
- Common in search engines, chatbots, and plagiarism detection

Summary Table

Metric	Data Type	Math Form	Use Case
Euclidean	Numeric	$\sqrt{\sum (x_i - y_i)^2}$	K-Means Clustering
Manhattan	Numeric	$\sum x_i - y_i $	Lasso Regression
Jaccard	Set/Binary	$\frac{ A \cap B }{ A \cup B }$	Recommender Systems
Minkowski	Numeric	$(\sum x_i - y_i ^p)^{1/p}$	Custom KNN
Cosine	Vector/High-dim	$\frac{A \cdot B}{\ A\ \ B\ }$	Text Similarity

Some More Distance Measures

What is Mahalanobis Distance?

- Measures distance between a point and a distribution
- Accounts for correlation between variables
- Scale-invariant
- More robust than Euclidean distance for multivariate data

- Euclidean distance treats all variables equally
- Mahalanobis considers variable correlation and scale
- Tells how many standard deviations a point is from the mean
- Adapts to the shape and orientation of the data

Mathematical Formulation

Given:

- Point: \mathbf{x}
- Mean: $\boldsymbol{\mu}$
- Covariance: Σ

$$D_M(x) = \sqrt{(\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu})}$$

- Σ^{-1} : Inverse of covariance matrix
- Output is a scalar distance

Properties

- Accounts for scale and correlation of features
- Reduces to Euclidean distance if $\Sigma = I$
- $D_M(x) = 0$: Point at mean
- $D_M(x) = 1$: One standard deviation away
- Useful for identifying multivariate outliers

- Outlier detection
- Clustering (e.g., GMMs)
- Multivariate hypothesis testing
- Pattern recognition (e.g., face recognition)
- Quality control in industrial settings

Example Calculation

- Mean: $\mu = \begin{bmatrix} 5 \\ 3 \end{bmatrix}$
- Covariance: $\Sigma = \begin{bmatrix} 4 & 2 \\ 2 & 3 \end{bmatrix}$
- Point: $\mathbf{x} = \begin{bmatrix} 6 \\ 4 \end{bmatrix}$

$$D_M(x) = \sqrt{(x - \mu)^T \Sigma^{-1} (x - \mu)}$$

Python Implementation

```
import numpy as np
from scipy.spatial import distance

data = np.array([[2, 3], [3, 5], [4, 8], [5, 11]])
x = np.array([3, 7])
mean = np.mean(data, axis=0)
cov = np.cov(data, rowvar=False)
inv_cov = np.linalg.inv(cov)

# Manual calculation
diff = x - mean
d_mahal = np.sqrt(diff.T @ inv_cov @ diff)

# Using scipy
d_scipy = distance.mahalanobis(x, mean, inv_cov)
```

Tips & Best Practices

- Ensure data follows approximately multivariate normal distribution
- Handle singular covariance matrices (regularization, PCA)
- Normalize data if using with other distance metrics
- Be cautious with high-dimensional data (curse of dimensionality)

What is MCC?

- Metric for evaluating binary (and multiclass) classification
- Considers: TP, TN, FP, FN
- Balanced measure, even with class imbalance
- Interpreted as a correlation coefficient

Intuition

- Accuracy fails in imbalanced datasets
- MCC reflects the relationship between predictions and ground truth
- Robust and symmetric

Values:

- 1 = Perfect prediction
- 0 = No better than random
- -1 = Total disagreement

MCC Formula

$$\text{MCC} = \frac{(TP \cdot TN) - (FP \cdot FN)}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

- If denominator = 0, MCC is defined as 0
- Normalized between -1 and 1

- $-1 \leq \text{MCC} \leq 1$
- Handles class imbalance
- Invariant to label flipping
- More reliable than accuracy or F1-score alone

- Medical diagnostics
- Fraud detection
- Binary classifiers in NLP, vision
- Model selection and benchmarking
- Extended to multiclass and multilabel problems

Example Calculation

Confusion Matrix:

- $TP = 70$, $TN = 50$
- $FP = 10$, $FN = 5$

$$\begin{aligned} MCC &= \frac{(70 \cdot 50) - (10 \cdot 5)}{\sqrt{(70 + 10)(70 + 5)(50 + 10)(50 + 5)}} \\ &= \frac{3450}{\sqrt{19800000}} \approx 0.776 \end{aligned}$$

```
from sklearn.metrics import matthews_corrcoef

y_true = [1, 1, 0, 1, 0, 0, 1]
y_pred = [1, 0, 0, 1, 0, 1, 1]

mcc = matthews_corrcoef(y_true, y_pred)
print("MCC:", mcc)
```

Manual Calculation

```
from math import sqrt

TP, TN, FP, FN = 70, 50, 10, 5
numerator = (TP * TN) - (FP * FN)
denominator = sqrt((TP + FP)*(TP + FN)*(TN + FP)*(TN + FN))
mcc = numerator / denominator if denominator != 0 else 0
print("MCC (manual):", mcc)
```

- Use MCC for imbalanced datasets
- Report MCC with other metrics (e.g., AUC, F1)
- Useful in production to track model drift
- Preferable when false positives/negatives carry different costs

Thank you!